

JW_Data_Pipelines

James Waterford

2023-03-07

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'  
  
## The following objects are masked from 'package:stats':  
##  
##   filter, lag  
  
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

There are two main ways to run a code:

Nested Code

This method involves running multiple codes in a single line.

```
numbers <- 1:300  
mean(numbers)
```

```
## [1] 150.5
```

```
sqrt(mean(numbers))
```

```
## [1] 12.26784
```

Sequential Code

This method generates intermediate variables to perform statistics on.

```
numbers <- -300:456  
mn <- mean(numbers)  
sqrt(mn)
```

```
## [1] 8.831761
```

```
library(readr)
surveys <- read_csv("197-raw_storage/surveys.csv")
```

```
## Rows: 35549 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr (2): species_id, sex
## dbl (7): record_id, month, day, year, plot_id, hindfoot_length, weight
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
species_data <- read.csv( "197-raw_storage/species.csv")
plots_data <- read.csv("197-raw_storage/plots.csv")
```

```
## # A tibble: 6 x 4
##   year month   day species_id
##   <dbl> <dbl> <dbl> <chr>
## 1  1977     7    16 NL
## 2  1977     7    16 NL
## 3  1977     7    16 DM
## 4  1977     7    16 DM
## 5  1977     7    16 DM
## 6  1977     7    16 PF
```

```
## # A tibble: 6 x 4
##   year species_id weight weight_kg
##   <dbl> <chr>      <dbl>      <dbl>
## 1  1977 PF         4        4000
## 2  1981 PF         4        4000
## 3  1981 PF         4        4000
## 4  1982 PF         4        4000
## 5  1982 PF         4        4000
## 6  1983 RM         4        4000
```

Pipe

Pipes can be implemented in R with the `dplyr` package, and the `magrittr` package.

The original symbol of the pipe is `%>%`. However, we can also use `|>` for the same effect. The purpose of this pipe is to eliminate or reduce the need of intermediate variables. R Studio includes a shortcut for the pipe: `cmd + shift + m`

```
library(magrittr)
1:300 |> mean() |> sqrt() -> mean_square
```

When we use a pipeline, we don't need to plug in the variable name every time. This was a good practice run, but let's load some real data now.

Let's calculate the median year of surveys.

```
library(readr)
surveys <- read_csv("197-raw_storage/surveys.csv")

## Rows: 35549 Columns: 9
## -- Column specification -----
## Delimiter: ","
## chr (2): species_id, sex
## dbl (7): record_id, month, day, year, plot_id, hindfoot_length, weight
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
surveys$year %>% median()
```

```
## [1] 1990
```

Let's try calculating the mean of the weight. Because there are NAs in our weight column, we'll need to remove these.

```
surveys$weight |> mean(na.rm=TRUE)
```

```
## [1] 42.67243
```

Data Manipulation Practice

Sometimes it is much easier to run keep editing a data set, until it matches your intentions.

```
surveys2 <- select(surveys, year, species_id, weight) |>
  mutate(weight_kg = weight/1000) |>
  filter(!is.na(weight_kg)) |>
  select(year, species_id, weight_kg)

str(surveys2)

## tibble [32,283 x 3] (S3: tbl_df/tbl/data.frame)
## $ year      : num [1:32283] 1977 1977 1977 1977 1977 ...
## $ species_id: chr [1:32283] "DM" "DM" "DM" "DM" ...
## $ weight_kg : num [1:32283] 0.04 0.048 0.029 0.046 0.036 0.052 0.008 0.022 0.035 0.007 ...

# surveys[ , c(1,3)]
# surveys[ , c("year", "weight_kg")]
```

Let's try one more example

The following code is written using intermediate variables. It obtains the data for "DS" in the "species_id" column, sorted by year, with only the year and weight columns. Write the same code to get the same output but using pipes instead.

```
ds_data <- filter(surveys, species_id == "DS", !is.na(weight))
ds_data_by_year <- arrange(ds_data, year)
ds_weight_by_year <- select(ds_data_by_year, year, weight)
```

```
filter(surveys, species_id == "DS", !is.na(weight)) |>
  arrange(year) |>
  select(year, weight) -> ds_data_by_year
head(ds_data_by_year)
```

```
## # A tibble: 6 x 2
##   year weight
##   <dbl> <dbl>
## 1  1977    117
## 2  1977    121
## 3  1977    115
## 4  1977    120
## 5  1977    118
## 6  1977    126
```

What if I want to pipe to an argument other than the first argument?

```
str(surveys)
```

```
## spc_tbl_ [35,549 x 9] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ record_id      : num [1:35549] 1 2 3 4 5 6 7 8 9 10 ...
## $ month          : num [1:35549] 7 7 7 7 7 7 7 7 7 7 ...
## $ day            : num [1:35549] 16 16 16 16 16 16 16 16 16 16 ...
## $ year           : num [1:35549] 1977 1977 1977 1977 1977 ...
## $ plot_id        : num [1:35549] 2 3 2 7 3 1 2 1 1 6 ...
## $ species_id     : chr [1:35549] "NL" "NL" "DM" "DM" ...
## $ sex            : chr [1:35549] "M" "M" "F" "M" ...
## $ hindfoot_length: num [1:35549] 32 33 37 36 35 14 NA 37 34 20 ...
## $ weight         : num [1:35549] NA NA NA NA NA NA NA NA NA ...
## - attr(*, "spec")=
## .. cols(
## ..   record_id = col_double(),
## ..   month = col_double(),
## ..   day = col_double(),
## ..   year = col_double(),
## ..   plot_id = col_double(),
## ..   species_id = col_character(),
## ..   sex = col_character(),
## ..   hindfoot_length = col_double(),
## ..   weight = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
lm(formula = weight ~ year, data = surveys)
```

```
##
## Call:
## lm(formula = weight ~ year, data = surveys)
##
## Coefficients:
```

```
## (Intercept)      year
##    2752.137      -1.361
```

Sometimes, us coders are lazy. We don't want to put in every variable detail if we can avoid it. So we use the pipeline.

```
surveys %>%
  lm(formula = weight ~ year, data = _)
##This code will not run because we called data incorrectly##
surveys %>%
  lm(formula = weight ~ year, data = .)
surveys |>
  lm(formula = weight ~ year, data = _)
```

Piping Placeholders

```
filter(surveys, species_id == "DS", !is.na(weight)) %>%
  lm(formula = weight ~ year, data = .) %>%
  summary()
```

```
##
## Call:
## lm(formula = weight ~ year, data = .)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -109.787  -12.440    3.723   14.886   69.886
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -709.1968   263.2510  -2.694  0.00711 **
## year         0.4184     0.1328    3.150  0.00165 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22.86 on 2342 degrees of freedom
## Multiple R-squared:  0.00422,    Adjusted R-squared:  0.003795
## F-statistic: 9.925 on 1 and 2342 DF,  p-value: 0.001651
```

Data Grouping / Data Aggregation

The function `group_by()` combines rows based on *matching columns*. `group_by([data], [column])`

```
group_by(surveys, year)
```

```
## # A tibble: 35,549 x 9
## # Groups:   year [26]
##   record_id month   day year plot_id species_id sex hindfoot_length weight
##         <dbl> <dbl> <dbl> <dbl>   <dbl>   <chr>      <chr>          <dbl>   <dbl>
```

```
## 1      1      7      16 1977      2 NL      M      32      NA
## 2      2      7      16 1977      3 NL      M      33      NA
## 3      3      7      16 1977      2 DM      F      37      NA
## 4      4      7      16 1977      7 DM      M      36      NA
## 5      5      7      16 1977      3 DM      M      35      NA
## 6      6      7      16 1977      1 PF      M      14      NA
## 7      7      7      16 1977      2 PE      F      NA      NA
## 8      8      7      16 1977      1 DM      M      37      NA
## 9      9      7      16 1977      1 DM      F      34      NA
## 10     10     7      16 1977      6 PF      F      20      NA
## # ... with 35,539 more rows
```

```
surveys %>%
  group_by(year)
```

```
## # A tibble: 35,549 x 9
## # Groups:   year [26]
##   record_id month   day  year plot_id species_id sex  hindfoot_length weight
##   <dbl> <dbl> <dbl> <dbl> <dbl> <chr>      <chr>      <dbl> <dbl>
## 1      1      7    16 1977      2 NL      M          32      NA
## 2      2      7    16 1977      3 NL      M          33      NA
## 3      3      7    16 1977      2 DM      F          37      NA
## 4      4      7    16 1977      7 DM      M          36      NA
## 5      5      7    16 1977      3 DM      M          35      NA
## 6      6      7    16 1977      1 PF      M          14      NA
## 7      7      7    16 1977      2 PE      F          NA      NA
## 8      8      7    16 1977      1 DM      M          37      NA
## 9      9      7    16 1977      1 DM      F          34      NA
## 10     10     7    16 1977      6 PF      F          20      NA
## # ... with 35,539 more rows
```

```
surveys %>%
  group_by(sex, year) %>%
  summarize()
```

```
## 'summarise()' has grouped output by 'sex'. You can override using the '.groups'
## argument.
```

```
## # A tibble: 78 x 2
## # Groups:   sex [3]
##   sex    year
##   <chr> <dbl>
## 1 F     1977
## 2 F     1978
## 3 F     1979
## 4 F     1980
## 5 F     1981
## 6 F     1982
## 7 F     1983
## 8 F     1984
## 9 F     1985
## 10 F    1986
## # ... with 68 more rows
```

Okay, this is an alright tool, but it's better when we know how to use it.

```
group_by(surveys, sex, year) %>%  
  summarize(count = n())
```

```
## 'summarise()' has grouped output by 'sex'. You can override using the '.groups'  
## argument.
```

```
## # A tibble: 78 x 3  
## # Groups:   sex [3]  
##   sex    year count  
##   <chr> <dbl> <int>  
## 1 F      1977   204  
## 2 F      1978   503  
## 3 F      1979   327  
## 4 F      1980   605  
## 5 F      1981   631  
## 6 F      1982   823  
## 7 F      1983   771  
## 8 F      1984   445  
## 9 F      1985   636  
## 10 F     1986   414  
## # ... with 68 more rows
```

```
group_by(surveys, sex, year) %>%  
  summarize(mean = mean(weight, na.rm = TRUE))
```

```
## 'summarise()' has grouped output by 'sex'. You can override using the '.groups'  
## argument.
```

```
## # A tibble: 78 x 3  
## # Groups:   sex [3]  
##   sex    year mean  
##   <chr> <dbl> <dbl>  
## 1 F      1977  47.6  
## 2 F      1978  70.0  
## 3 F      1979  65.6  
## 4 F      1980  57.4  
## 5 F      1981  63.4  
## 6 F      1982  55.4  
## 7 F      1983  55.9  
## 8 F      1984  49.0  
## 9 F      1985  47.1  
## 10 F     1986  54.7  
## # ... with 68 more rows
```

```
surveys %>%  
  group_by(species_id) %>%  
  summarize(count = n())
```

```
## # A tibble: 49 x 2
```

```
##   species_id count
##   <chr>      <int>
## 1 AB         303
## 2 AH         437
## 3 AS          2
## 4 BA         46
## 5 CB         50
## 6 CM         13
## 7 CQ         16
## 8 CS          1
## 9 CT          1
## 10 CU         1
## # ... with 39 more rows
```

```
surveys %>%
  group_by(species_id, year) %>%
  summarize(count = n())
```

'summarise()' has grouped output by 'species_id'. You can override using the
'.groups' argument.

```
## # A tibble: 535 x 3
## # Groups:   species_id [49]
##   species_id year count
##   <chr>      <dbl> <int>
## 1 AB         1980     5
## 2 AB         1981     7
## 3 AB         1982    34
## 4 AB         1983    41
## 5 AB         1984    12
## 6 AB         1985    14
## 7 AB         1986     5
## 8 AB         1987    35
## 9 AB         1988    39
## 10 AB        1989    31
## # ... with 525 more rows
```

```
surveys %>%
  filter(species_id == "DO") %>%
  group_by(year) %>%
  summarize(mean = mean(weight, na.rm = TRUE))
```

```
## # A tibble: 26 x 2
##   year mean
##   <dbl> <dbl>
## 1 1977 42.7
## 2 1978 45
## 3 1979 45.9
## 4 1980 48.1
## 5 1981 49.1
## 6 1982 47.9
## 7 1983 47.2
```



```
## 8 1984 48.4
## 9 1985 48.0
## 10 1986 49.4
## # ... with 16 more rows
```